



CITS5508 Machine Learning Semester 1, 2022

Lab Sheet 5

Assessed, worth 20%. Due: 11:59pm, Friday 27th May 2022

1 Outline

This lab sheet consists of a large project using the CIFAR-10 dataset. This is a data set containing 10 classes of colour images of size 32×32 pixels. The training set is perfectly balanced, with 6,000 images per class. The test set contains 10,000 instances. Your task for this labsheet is to train an MLP and a CNN for the classification task and compare their performance. This lab sheet is a good practical exercise to test your understanding of the techniques covered in Chapters 10, 11, and 14.

2 Submission

Name your Jupyter Notebook file as **Surname_FirstName-lab05.ipynb** (please replace **Surname** and **FirstName**¹ by your own surname and first name). For your trained MLP and CNN models, name them as **Surname_FirstName-MLP** and **Surname_FirstName-CNN**. Your models are likely to be saved as a small number of files so **Surname_FirstName-MLP** and **Surname_FirstName-CNN** would be subdirectories containing those files. You need to make sure that these subdirectories were created in the same directory with your Notebook file, as this is what we will assume in our marking process.

NOTE 1: You should zip² your Notebook file and the two subdirectories together and submit **a single lab05.zip file**. When you create your zip file for submission, ensure that

- **no data files** are included. The dataset for this labsheet is very large and we do not need you to supply the dataset back to us.
- **no checkpoint files** are included. You can create these files to monitor the training progress of your models but these checkpoint files should be for yourself only. You should submit only the final trained models, as including all the checkpoint files will significantly increase the size of your zip file.
- **no hidden files/directories** are included. Jupyter-Lab and Jupyter-notebook create backup files in the **.ipynb_checkpoints** directory. This is a hidden directory that might not show up (on Windows, it will depend on whether you turn on the option of showing hidden files/directories; in Linux/MacOS, you can type: `ls -al` to list the contents (including all hidden files) of the current directory). On the MacOS, there are further hidden files/directories starting with dot-underscore (`._`) characters. None of these hidden files should be included in your submission.

Any of the above files, if found in your submission, will incur a 5% penalty to your mark.

NOTE 2: You must submit your zip file to **cssubmit**. You should reserve at least half an hour for the submission procedure as csmarks carries out virus checking (which can't be disabled) on each submitted

¹You can include your middle name if you like. It does not matter. We just want to be able to distinguish students' files when we unzip the whole class's submissions together in one directory. **NOTE:** If your surname and/or given name contain special characters, such as apostrophes, `'`, etc, you may have to omit them or replace them by underscores or hyphens as these characters have special meanings in Python: apostrophes are considered as single quotes; `'` is used to denote directories.

²On Linux and MacOS, type in a terminal window:

```
zip -r lab05.zip Surname_FirstName-lab05.ipynb Surname_FirstName-MLP Surname_FirstName-CNN
```

On Windows, see examples on the web, e.g.,

<https://support.microsoft.com/en-us/windows/zip-and-unzip-files-8d28fa72-f2f9-712f-67df-f80cf89fd4e5>

file. Your zip file for the lab sheet can be quite large as it needs to include two trained models. You need to ensure that your internet connection is stable during the entire uploading process. If possible, have your computer wire-connected to your modem before you start your submission.

You should wait until you get an acknowledgement screen confirming that your submission has completed successfully. You can take a screen shot and keep the image as evidence of your submission. Do **not** send your zip file via email to the Unit Coordinator. Do **not** upload your zip file onto Google Drive or OneDrive and then send the link to the Unit Coordinator. **Only files uploaded on csubmit will be marked.**

NOTE 3: Do not blindly copy examples that you found on the internet. Many of these examples are too complex and can only be trained on huge datasets such as ImageNet. These complex networks may make your zip file too large to be uploaded onto csubmit. **Significant penalty** will be imposed on your mark if your MLP and CNN are too complex and do not match the specification stated in the lab sheet. You should be able to complete this lab sheet by following examples in the lectures, textbook, and sample codes from the author.

3 Data Download and Preparation

CIFAR-10 can be downloaded from the website: <https://www.cs.toronto.edu/~kriz/cifar.html> which covers very clear description about the dataset. You need to download the **cifar-10-python.tar.gz** file from the link **CIFAR-10 python version**. The direct link is

<https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>

On Linux, you can unzip the file using the command `tar xvfz cifar-10-python.tar.gz`. On Windows, you may need to unzip twice, firstly to get the **cifar-10-python.tar** file and then unzip to get the data files. If your unzipping operation is successful, you should see a subdirectory named *cifar-10-batches-py*. Under there, you should find the following files:

```
batches.meta
data_batch_1
data_batch_2
data_batch_3
data_batch_4
data_batch_5
readme.html
test_batch
```

You should move all these files to the same directory with your **Surname_FirstName-lab05.ipynb** file. Here, the dataset has been organised as follows: The training set is split into five *data_batch* files. The test set is in the *test_batch* file. To help you read these files correctly, two additional files are supplied with this lab sheet:

- **data_loader.py**. This file contains a Python class called `DataLoader`. The function `load_batch` can be used to read the dataset.
- **lab05-sample.ipynb**. This file shows you how to call the `load_batch` function mentioned above. Feel free to modify this file to form your `Surname_FirstName-lab05.ipynb` file.

Note that the image pixel values have been normalised to the range 0..1 by the `load_batch` function. You shouldn't need to perform any further feature scaling.

4 Tasks

- (i) Use an 85/15 random split on the training set to form a validation set. Use these training, validation, and test sets for training, validating, and testing both of your MLP and CNN models later on.
- (ii) Write a small function that takes in appropriate arguments so that it can be used to display 20 randomly sampled images from the arguments. This function should be called 3 times – for the training, validation, and test sets. The figure displayed by the function should show the corresponding class name of each image.

(iii) Implementation of an MLP

Design an MLP that has 2 to 3 hidden layers and an output layer. Use an appropriate function from `tensorflow.keras` to show a summary of your MLP architecture.

You should try to design your MLP in such a way that **the hidden layers have fewer neurons than the input layer**. This would help to avoid overfitting. You should train the network for 100 epochs but use the validation set for *early stopping* (which also helps to overcome overfitting).

In each hidden layer, use a suitable number of neurons and an appropriate activation function. Experiment with **two** possible settings for each of the following hyperparameters:

- connection weight initialisation;
- learning rate scheduling (you will need to use *callback* and write a small function for this);
- dropout rate.

You can use grid search together with the validation set created above to find the optimal value for each of them. In the final version of your Notebook file, the code for this hyperparameter fine tuning process should be moved to a markdown cell and explained there.

Your trained MLP model should be saved to the directory named Surname_FirstName-MLP.

(iv) Implementation of a CNN

Design a CNN that has **2 to 3 convolutional layers** and a pooling layer between consecutive **convolutional layers**. Before the output layer, you will need to have 1 to 2 fully connected layers and maybe *batch normalisation layers* to **help control the numerical values of the network weights**. Use an appropriate function from `tensorflow.keras` to show a summary of your CNN architecture.

Similar to the MLP model above, you should explore **two** possible settings for each of the following hyperparameters:

- kernel size;
- number of kernels³;
- activation function.

For your CNN, you can use the optimal way to initialise the network, the **optimal learning rate scheduling, and the optimal dropout value** that you found from the training of your MLP model above.

Same as before, once you have got the optimal values of these hyperparameters, move the code into a markdown cell.

Your trained CNN model should be saved to the directory named Surname_FirstName-CNN.

(v) Structure of your code

In the marking process, we want to have the following options in your implementation of both the MLP and CNN:

³the term “kernel” and “filter” are often used interchangeably in computer vision

- (a) train each of your networks from scratch for 100 epochs using the optimal hyperparameter values that you have found; or
- (b) load each of your trained models and train it for just 1 more epoch.

After that, the model (obtained from a) or b)) should be used to do predictions on the training set and the test set.

So your Python code should look like this:

```
if the MLP model subdirectory is present in the current directory
    load the model,
    display its architecture,
    train for one epoch.
else
    set up the model and display its architecture,
    train the model from scratch for 100 epochs using the optimal
hyperparameter values.
    save the model to the directory described earlier.

use the model to perform predictions.
```

You should report and compare your trained model's classification accuracies, F1 scores, and show the confusion matrices on both the training and test sets.

Similar to the code for your MLP model, repeat the above steps for your CNN model.

- (vi) Compare and comment on your MLP and CNN models on the test set, in terms of: classification performance (accuracy, F1 score, precision per class), model complexity (e.g., number of trainable parameters), and computation time⁴.

For each model, display also a few correctly classified images and a few failure cases for the test set.

5 Google Colab

The training of your MLP and CNN will take a long time to complete if your computer does not have a GPU. It is very easy to upload all the batch files for the dataset, data_loader.py, and your Surname_FirstName-lab05.ipynb file onto your Google Drive. If you decide to use Google Colab, then you should find out how to mount the data files on your Google drive so that they are visible in your code.

By default, GPU is not available on Colab. To specify that you need a GPU, select the menu item *Runtime* and then the option *Change runtime type*. In the popped-up window, select "GPU" for *Hardware accelerator*.

6 High performance computing (HPC) facilities at UWA

UWA has a number of multi-core computers, some of which have GPUs. It is not compulsory to use these facilities for the labsheet. However, they will certainly help you cut down the training time of your deep networks. An account has already been set up for each student on **Kaya** (the name of the GPU host). Refer to the instructions given in the file **Using-Kaya.pdf** and watch the related video (available on LMS) if you want to use the UWA HPC facilities. The video was recorded in 2021, so there will be small differences from what you need to do this year, e.g., the Python environment name is not the same this year.

⁴You can use any of the Python modules, such as `time`, `timeit`, to help you.

7 Penalty on late submissions

See the URL below about late submission of assignments:

https://ipoint.uwa.edu.au/app/answers/detail/a_id/2711/~consequences-for-late-assignment-submission

Late submissions will be calculated as 5% of the total mark per 24 overdue hours. So if your submission is late by 3 hours, the penalty will only be 0.625%.

8 Plagiarism

You should attempt the lab sheet by yourself. Collusion with other students is considered to be serious academic misconduct and can cause you to be suspended or expelled from the unit. Please see <https://www.uwa.edu.au/students/my-course/student-conduct> for more details.